

HW3_Problem4-3

May 5, 2020

```
[0]: import json
import pandas as pd
import numpy as np
import time
import math
import string, re
import matplotlib.pyplot as plt
import sys
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from scipy.sparse import csr_matrix
from sklearn.svm import LinearSVC
from sklearn.linear_model import SGDClassifier
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[0]: f = open('/content/drive/My Drive/HW3/labeled_small.json','r')
data = json.load(f)
```

```
[0]: beer = pd.DataFrame.from_dict(data)
```

```
[5]: type(beer)
```

```
[5]: pandas.core.frame.DataFrame
```

```
[6]: beer
```

```
[6]:
```

	style	palate	...	beer_name	beer_id
0	17	3.0	...	John Harvards Simcoe IPA	45842
1	17	4.0	...	John Harvards Simcoe IPA	45842
2	33	3.0	...	John Harvards Cristal Pilsner	95213
3	33	2.0	...	John Harvards Fancy Lawnmower Beer	65957
4	58	4.0	...	John Harvards Vanilla Black Velvet Stout	41336

...
99995	69	2.0	...	Mission Street Pale Ale	68983	
99996	69	2.0	...	Mission Street Pale Ale	68983	
99997	69	3.0	...	Mission Street Pale Ale	68983	
99998	69	3.0	...	Mission Street Pale Ale	68983	
99999	69	2.0	...	Mission Street Pale Ale	68983	

[100000 rows x 11 columns]

```
[0]: with open('/content/drive/My Drive/HW3/vocab_30.json', 'r') as f:
      vocab = json.load(f)
```

```
[8]: len(vocab)
```

```
[8]: 41346
```

```
[0]: # Part 1: Data Inspection
```

```
[10]: beer.describe()
```

```
[10]:
```

	style	palate	...	aroma	beer_id
count	100000.000000	100000.000000	...	100000.000000	100000.000000
mean	36.464290	3.277750	...	6.447190	55650.48179
std	22.538597	0.773156	...	1.439254	33327.57236
min	0.000000	1.000000	...	1.000000	1.000000
25%	16.000000	3.000000	...	6.000000	26244.00000
50%	37.000000	3.000000	...	7.000000	56100.00000
75%	58.000000	4.000000	...	7.000000	86041.00000
max	73.000000	5.000000	...	10.000000	110283.00000

[8 rows x 9 columns]

```
[0]: # Mean: 13.396
      # Median: 14.000
      # Standard Deviation: 2.949
```

```
[0]: # Now, do people have a similar taste? In order to answer this question,
      # We can assume that if people do, they will leave similar ratings for
      # the same kind of beers. Then, the standard deviation of ratings for
      # a selected beer should be quite small compared to the overall ratings
      # for all kinds of beers
```

```
[0]: # Let's try three examples
```

```
[0]: beer1 = beer[beer["beer_name"] == "Mission Street Pale Ale"]
      beer2 = beer[beer["beer_name"] == "Burggraf Maibock"]
      beer3 = beer[beer["beer_name"] == "Barley Island Barrel-Aged Count Hopula"]
```

```
[15]: beer1["overall"].std()
```

```
[15]: 2.1420781684525454
```

```
[16]: beer2["overall"].std()
```

```
[16]: 1.789114825239473
```

```
[17]: beer3["overall"].std()
```

```
[17]: 0.8997354108424372
```

```
[0]: # The standard deviation of the overall rating of all beers was 2.949.  
# We can see that the overall ratings of the above three beers have standard  
# deviations lower than this threshold, but some are much lower than others.  
# We can argue that people could have similar tastes for only some types of  
→beer
```

```
[0]: # Part 2
```

```
[0]: # (a) Generating Features
```

```
[0]: # Splitting data into training, validation, and testing
```

```
[0]: # But First, create a function that gives binary labels
```

```
[0]: label = np.array(beer["overall"] >= 14).astype(int)
```

```
[0]: train_data, test_data, train_label, test_label = train_test_split(beer, label,  
→test_size = 0.3, random_state =2)  
val_data, test_data, val_label, test_label = train_test_split(test_data,  
→test_label, test_size = 0.5, random_state =2)
```

```
[0]: # function that normalize the text
```

```
[0]: vocab_index = {}  
vocab_lst = list(vocab.keys())  
for i, vocab in enumerate(vocab_lst):  
    vocab_index[vocab] = i
```

```
[0]: def normalize_txt(data):  
    r = re.compile("[ " + re.escape(string.punctuation) + " ]")  
    normal = []  
    for i, row in data.iterrows():  
        words = r.sub('', row['review'])  
        split = words.split(' ')  
        lst = []
```

```

        for s in split:
            if s in vocab_index:
                lst.append(s)
        normal.append(lst)
    return normal

```

```
[0]: # function that creates sparse matrix
```

```
[0]: def create_sparse(document, vocab_index_):
    indptr = [0]
    indices = []
    data = []
    for doc in document:
        terms = set(doc)
        for term in terms:
            index = vocab_index_[term]
            indices.append(index)
            data.append(1)
        indptr.append(len(indices))
    sparse = csr_matrix((data, indices, indptr), dtype = bool).astype(int)
    return sparse

```

```
[30]: time_start = time.time()
norm_doc = normalize_txt(train_data)
print(time.time() - time_start)
```

9.474142074584961

```
[31]: time_start = time.time()
sparse = create_sparse(norm_doc, vocab_index)
print(time.time() - time_start)
```

1.0818650722503662

```
[0]: # It takes about 1.117 seconds to go over the entire training data set, which
    ↪ is
    # 70,000 documents long
```

```
[0]: # part (b): Logistic Regression
```

```
[0]: sparse_test = create_sparse(normalize_txt(test_data), vocab_index)
```

```
[0]: sparse_val = create_sparse(normalize_txt(val_data), vocab_index)
```

```
[36]: start = time.time()
lg = LogisticRegression(penalty = "l2", C = 1/10, solver = 'newton-cg')
lg.fit(sparse, train_label)
```

```
print(time.time() - start)
```

3.860454559326172

```
[37]: score = accuracy_score(val_label, lg.predict(sparse_val))
      score
```

[37]: 0.7605333333333333

```
[0]: # selecting the optimal lambda_ by trying different values on the validation set
```

```
[39]: lambda_set1 = [1, 5, 10, 20, 30, 40, 50, 70, 100, 300, 500, 1000]
      accuracy1 = []
      for lam in lambda_set1:
          print("lambda: {} | Penalty C: {}".format(lam, 1/lam))
          lg = LogisticRegression(penalty = "l2", C = 1/lam, solver = 'newton-cg')
          start = time.time()
          lg.fit(sparse, train_label)
          score = accuracy_score(val_label, lg.predict(sparse_val))
          accuracy1.append(score)
          print("operation time: %s" %(time.time() - start))
          print("accuracy: %s" %(score))
```

```
lambda: 1 | Penalty C: 1.0
operation time: 7.5562450885772705
accuracy: 0.7518
lambda: 5 | Penalty C: 0.2
operation time: 4.4831624031066895
accuracy: 0.7587333333333334
lambda: 10 | Penalty C: 0.1
operation time: 3.943840503692627
accuracy: 0.7605333333333333
lambda: 20 | Penalty C: 0.05
operation time: 3.605893135070801
accuracy: 0.7625333333333333
lambda: 30 | Penalty C: 0.03333333333333333
operation time: 3.516540288925171
accuracy: 0.7606666666666667
lambda: 40 | Penalty C: 0.025
operation time: 3.1036739349365234
accuracy: 0.7603333333333333
lambda: 50 | Penalty C: 0.02
operation time: 2.9787139892578125
accuracy: 0.7588
lambda: 70 | Penalty C: 0.014285714285714285
operation time: 2.881065607070923
accuracy: 0.7563333333333333
lambda: 100 | Penalty C: 0.01
```


```

operation time: 2.623732566833496
accuracy: 0.7542
lambda: 300 | Penalty C: 0.003333333333333335
operation time: 2.144583225250244
accuracy: 0.744
lambda: 500 | Penalty C: 0.002
operation time: 2.358748435974121
accuracy: 0.7391333333333333
lambda: 1000 | Penalty C: 0.001
operation time: 1.7104160785675049
accuracy: 0.7284666666666667

```

```
[40]: plt.plot(lambda_set1, accuracy1, 'bo')
```

```
[40]: [<matplotlib.lines.Line2D at 0x7f0ad94d3f28>]
```



```
[0]: # The accuracy rate was the highest when lambda = 20.
      # With lambda = 20, it took about 3.5 seconds to run the algorithm in order to
      # attain accuracy rate of 0.7625
```

```
[0]: # Moving on to the LinearSVC class
```

```
[43]: lambda_set2 = [1, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100]
      accuracy2 = []
      for lam in lambda_set2:
          print("lambda: {} | penalty C: {}".format(lam, 1 / lam))
```

```

lsvc = LinearSVC(penalty = 'l2', loss = 'hinge', C = 1/lam)
start = time.time()
lsvc.fit(sparse, train_label)
score = accuracy_score(val_label, lsvc.predict(sparse_val))
accuracy2.append(score)
print("operation time: %s " % (time.time() - start))
print("accuracy: %s"%(score))

```

lambda: 1 | penalty C: 1.0

/usr/local/lib/python3.6/dist-packages/sklearn/svm/_base.py:947:

ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

"the number of iterations.", ConvergenceWarning)

operation time: 5.048051595687866

accuracy: 0.7420666666666667

lambda: 10 | penalty C: 0.1

operation time: 1.37498140335083

accuracy: 0.7575333333333333

lambda: 20 | penalty C: 0.05

operation time: 0.6204280853271484

accuracy: 0.7577333333333334

lambda: 30 | penalty C: 0.03333333333333333

operation time: 0.5432126522064209

accuracy: 0.7586

lambda: 40 | penalty C: 0.025

operation time: 0.4545018672943115

accuracy: 0.7587333333333334

lambda: 50 | penalty C: 0.02

operation time: 0.33524274826049805

accuracy: 0.759

lambda: 60 | penalty C: 0.016666666666666666

operation time: 0.3623623847961426

accuracy: 0.7580666666666667

lambda: 70 | penalty C: 0.014285714285714285

operation time: 0.28708910942077637

accuracy: 0.7578666666666667

lambda: 80 | penalty C: 0.0125

operation time: 0.25193142890930176

accuracy: 0.7568666666666667

lambda: 90 | penalty C: 0.011111111111111112

operation time: 0.23581838607788086

accuracy: 0.7565333333333333


lambda: 100 | penalty C: 0.01

operation time: 0.2680075168609619

accuracy: 0.7561333333333333

```
[44]: plt.plot(lambda_set2, accuracy2, 'ro')
```

```
[44]: [<matplotlib.lines.Line2D at 0x7f0ad9452550>]
```



```
[0]: # Using hinge loss instead of logistic loss, lambda = 50 seems to be close to
# the optimal lambda value. It provided an accuracy rate of 0.759 with
# operation time of about 0.26 seconds. We can conclude that under the given
# circumstances, the logistic Regression model showed slightly higher (very
→small
# difference) accuracy, but the lsvc model computed the algorithms faster.
```

```
[0]: # part (c): Stochastic gradient descent
```

```
[0]: def yield_batches(x, y, epochs):
    np.random.shuffle(x)
    for e in range(epochs):
        interval = x.shape[0] // epochs
        yield x[int(e*interval): int((e + 1)*interval)], np.
→array(y[int(e*interval): int((e + 1)*interval)])[:, np.newaxis]
```

```
[0]: def compute_h(x, coef):
    return 1 / (1 + np.exp(-np.dot(x, coef)))
```

```
[0]: def compute_gradient(x, y, coef):
    h = compute_h(x, coef)
```




```
return (h - y) * x
```

```
[0]: def compute_error_rate(x, y, coef):  
    pred = np.matmul(x, coef)  
    y = np.array(y)[: , np.newaxis]  
    return np.mean(1/2 * (np.power((pred - y), 2)))
```

```
[0]: def perform_sgd(x, y, learning_rate, lambda_, epochs, decrease = False, plot =   
    ↪ True):  
    if type(x) != np.ndarray:  
        x = x.toarray()  
    x = np.concatenate((np.ones(len(x))[: , np.newaxis], x), axis = 1)  
    coef = np.zeros((x.shape[1], 1))  
    batches = yield_batches(x, y, epochs)  
    er = []  
    for i in range(epochs):  
        x_b , y_b = next(batches)  
        gradient = np.zeros((len(coef), 1))  
        error = compute_error_rate(x, y, coef)  
        er.append(error)  
        for j in range(len(x_b)):  
            gradient += (compute_gradient(x_b[j], y_b[j], coef)[: , np.newaxis])  
        if decrease:  
            coef = coef - learning_rate/(i+1) * (gradient / len(x_b) + lambda_  
    ↪ * coef)  
        else:  
            coef = coef - learning_rate * (gradient / len(x_b) + lambda_ * coef)  
  
    if plot:  
        plt.plot(er, "g")  
  
    return coef
```


```
[0]: # Performing Stochastic Gradient Descent without decreasing learning rate
```

```
[53]: coef = perform_sgd(sparse[:10000], train_label[:10000], 0.01, 0.1, 100,  
    ↪ decrease = False, plot = True)
```



```
[0]: # Performing Stochastic Gradient Descent with decreasing learning rate
```

```
[55]: coef = perform_sgd(sparse[:10000], train_label[:10000], 0.01, 0.1, 100, ↵
      ↪ decrease = True, plot = True)
```



```
[0]: # We can see that the error rate falls and converges as iterations proceed.
# When we use decreasing learning rate, we can erase most of the noise that
# causes oscillation and reach convergence much faster
```

```
[0]: # compute accuracy of the model by adopting misclassification threshold as 0.5
```

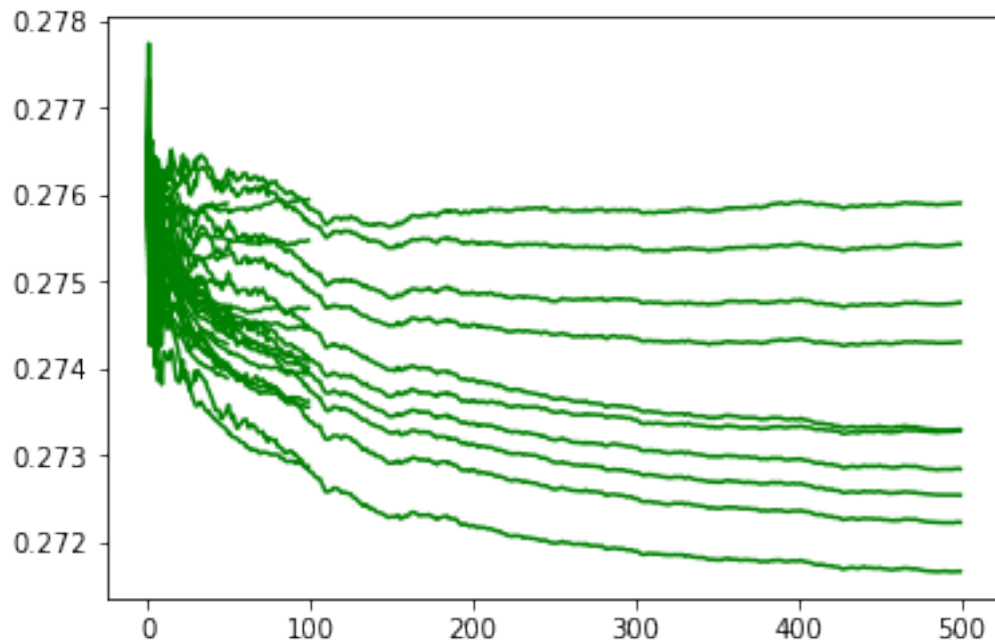
```
[0]: def compute_mis(coef, x_n, y_n):
    x = np.concatenate((np.ones(len(x_n))[:,np.newaxis], x_n), axis = 1)
    eta = np.matmul(x, coef)
    c = 0
    for i in range(len(y_n)):
        if eta[i] > 0:
            if y_n[i] == 0:
                c += 1
        elif eta[i] < 0 :
            if y_n[i] == 1:
                c += 1
    return c / len(y_n)
```

```
[0]: # now, let's test with different lambdas and number of iterations
```

```
[60]: s = sparse_val.toarray()
lambda_s = [0.0001, 0.001, 0.01, 0.1, 1, 10, 20, 30, 50, 100]
epochs = [10, 50, 100, 500]
for lam in lambda_s:
    for epoch in epochs:
        coef = perform_sgd(sparse[: 5000], train_label[: 5000], 0.01, lam,
→epoch, decrease = True, plot = True)
        error = compute_mis(coef, s, val_label)
        print("lambda: {} | epoch: {} | error rate: {}".format(lam, epoch,
→error))
```

```
lambda: 0.0001 | epoch: 10 | error rate: 0.4554
lambda: 0.0001 | epoch: 50 | error rate: 0.4554
lambda: 0.0001 | epoch: 100 | error rate: 0.45546666666666667
lambda: 0.0001 | epoch: 500 | error rate: 0.45586666666666664
lambda: 0.001 | epoch: 10 | error rate: 0.4554
lambda: 0.001 | epoch: 50 | error rate: 0.4554
lambda: 0.001 | epoch: 100 | error rate: 0.4554
lambda: 0.001 | epoch: 500 | error rate: 0.4554
lambda: 0.01 | epoch: 10 | error rate: 0.4554
lambda: 0.01 | epoch: 50 | error rate: 0.4554
lambda: 0.01 | epoch: 100 | error rate: 0.4554
lambda: 0.01 | epoch: 500 | error rate: 0.4554
lambda: 0.1 | epoch: 10 | error rate: 0.4554
lambda: 0.1 | epoch: 50 | error rate: 0.4554
lambda: 0.1 | epoch: 100 | error rate: 0.4554
```

lambda: 0.1 | epoch: 500 | error rate: 0.4552
lambda: 1 | epoch: 10 | error rate: 0.4554
lambda: 1 | epoch: 50 | error rate: 0.4554
lambda: 1 | epoch: 100 | error rate: 0.4554
lambda: 1 | epoch: 500 | error rate: 0.4554
lambda: 10 | epoch: 10 | error rate: 0.4554
lambda: 10 | epoch: 50 | error rate: 0.4554
lambda: 10 | epoch: 100 | error rate: 0.4554
lambda: 10 | epoch: 500 | error rate: 0.4554
lambda: 20 | epoch: 10 | error rate: 0.4554
lambda: 20 | epoch: 50 | error rate: 0.4554
lambda: 20 | epoch: 100 | error rate: 0.4554
lambda: 20 | epoch: 500 | error rate: 0.4554
lambda: 30 | epoch: 10 | error rate: 0.4554
lambda: 30 | epoch: 50 | error rate: 0.4554
lambda: 30 | epoch: 100 | error rate: 0.4554
lambda: 30 | epoch: 500 | error rate: 0.4554
lambda: 50 | epoch: 10 | error rate: 0.4554
lambda: 50 | epoch: 50 | error rate: 0.4554
lambda: 50 | epoch: 100 | error rate: 0.4554
lambda: 50 | epoch: 500 | error rate: 0.4554
lambda: 100 | epoch: 10 | error rate: 0.4554
lambda: 100 | epoch: 50 | error rate: 0.4554
lambda: 100 | epoch: 100 | error rate: 0.4554
lambda: 100 | epoch: 500 | error rate: 0.4554



```
[0]: # We can observe that the accuracy rate decreases as the size of batches
# increase. The optimal value of lambda seems to be between 10 and 30.
# We can narrow down the scope of optimal lambdas and number of iterations.
```


```
[65]: lambda_s = [0.01, 0.05, 0.1, 1, 10, 15, 20, 25, 30]
epochs = [250, 500, 750, 1000]
error_rate = []
lambda_ = []
epoch_ = []
for epoch in epochs:
    for lam in lambda_s:
        coef = perform_sgd(sparse[:5000], train_label[:5000], 0.01, lam, epoch,
        decrease = True)
        error = compute_mis(coef, s, val_label)
        print("lambda: {} | epoch: {} | error: {}".format(lam, epoch, error))
        error_rate.append(error)
```

```
lambda: 0.01 | epoch: 250 | error: 0.4554
lambda: 0.05 | epoch: 250 | error: 0.4554
lambda: 0.1 | epoch: 250 | error: 0.4554
lambda: 1 | epoch: 250 | error: 0.4553333333333333
lambda: 10 | epoch: 250 | error: 0.4554
lambda: 15 | epoch: 250 | error: 0.4554
lambda: 20 | epoch: 250 | error: 0.4554
lambda: 25 | epoch: 250 | error: 0.4554
lambda: 30 | epoch: 250 | error: 0.4554
lambda: 0.01 | epoch: 500 | error: 0.45653333333333335
lambda: 0.05 | epoch: 500 | error: 0.45626666666666665
lambda: 0.1 | epoch: 500 | error: 0.45573333333333333
lambda: 1 | epoch: 500 | error: 0.45953333333333335
lambda: 10 | epoch: 500 | error: 0.4554
lambda: 15 | epoch: 500 | error: 0.4554
lambda: 20 | epoch: 500 | error: 0.4554
lambda: 25 | epoch: 500 | error: 0.4554
lambda: 30 | epoch: 500 | error: 0.4554
lambda: 0.01 | epoch: 750 | error: 0.47673333333333334
lambda: 0.05 | epoch: 750 | error: 0.46166666666666667
lambda: 0.1 | epoch: 750 | error: 0.4554
lambda: 1 | epoch: 750 | error: 0.45453333333333334
lambda: 10 | epoch: 750 | error: 0.4554
lambda: 15 | epoch: 750 | error: 0.45546666666666667
lambda: 20 | epoch: 750 | error: 0.4554
lambda: 25 | epoch: 750 | error: 0.4554
lambda: 30 | epoch: 750 | error: 0.45593333333333336
lambda: 0.01 | epoch: 1000 | error: 0.45553333333333335
lambda: 0.05 | epoch: 1000 | error: 0.45613333333333334
lambda: 0.1 | epoch: 1000 | error: 0.456
```

```

lambda: 1 | epoch: 1000 | error: 0.46
lambda: 10 | epoch: 1000 | error: 0.45553333333333335
lambda: 15 | epoch: 1000 | error: 0.4554
lambda: 20 | epoch: 1000 | error: 0.45546666666666667
lambda: 25 | epoch: 1000 | error: 0.4554
lambda: 30 | epoch: 1000 | error: 0.4554

```



```


[67]: plt.subplot(121)
plt.plot(epochs * len(lambda_s) , error_rate,'bo')
plt.xlabel("epochs")
plt.subplot(122)
plt.plot(lambda_s * len(epochs), error_rate,'go')
plt.xlabel("lambdas")

```

```

[67]: Text(0.5, 0, 'lambdas')

```



```
[0]: # The results show that the optimal lambda value and the number of epochs are
# about 1 and 400 respectively(based on the sample of training data we used).
# We can use these figures to model the test data
# and compare the results with the Logistic Regression function
```

```
[0]: st = sparse_test.toarray()
coef = perform_sgd(sparse[: 5000], train_label[: 5000], 0.01, 1, 750, decrease_
↪ = True, plot = False)
error = compute_mis(coef, st, test_label)
```

```
[69]: print("test error: ", error)
```

test error: 0.45053333333333334

```
[0]: # Although the speed of computation is much faster for the stoichastic
# gradient descent method, the logistic regression function in the previous
# item was slightly more accurate. When we implemented stoichastic gradient
# descent, the lowest error rate we could achieve on the test set was
# only 0.4497
```

```
[0]:
```

```
[0]: # Part 3: Scores versus test
```


```
[0]: # Function that compiles the scores for other criteria
```

```
[0]: def compile_scores(data):
    lst = []
    for i in range(len(data)):
        lst.append([data.iloc[i]['appearance'], data.iloc[i]['aroma'],\
                    data.iloc[i]['palate'], data.iloc[i]['style'], data.iloc[i]['taste']])
    return np.array(lst)
```

```
[0]: train_scores = compile_scores(train_data)
```

```
[0]: test_scores = compile_scores(test_data)
val_scores = compile_scores(val_data)
```

```
[76]: coef = perform_sgd(val_scores[:5000], val_label[:5000], 0.01, 10, 500, decrease_
    ↪ = True, plot = True)
```



```
[0]: # The results show that the loss function converges much quicker than the model
# using only overall ratings
```

```
[78]: lambda_s = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
epochs = [50, 100, 500]
error = []
epoch_ = []
lambda_ = []
for epoch in epochs:
    for lam in lambda_s:
```



```

coef = perform_sgd(train_scores[:5000], train_label[:5000], 0.01, lam, \
                    epoch, decrease = True, plot = False)
e = compute_mis(coef, val_scores, val_label)
print("lambda: {} | epoch: {} | error: {}".format(lam, epoch, e))
epoch_.append(epoch)
lambda_.append(lam)
error.append(e)

```

```

lambda: 0.001 | epoch: 50 | error: 0.4554
lambda: 0.01 | epoch: 50 | error: 0.4554
lambda: 0.1 | epoch: 50 | error: 0.4554
lambda: 1 | epoch: 50 | error: 0.4554
lambda: 10 | epoch: 50 | error: 0.4554
lambda: 100 | epoch: 50 | error: 0.4554
lambda: 1000 | epoch: 50 | error: 0.4612
lambda: 0.001 | epoch: 100 | error: 0.4554
lambda: 0.01 | epoch: 100 | error: 0.4554
lambda: 0.1 | epoch: 100 | error: 0.4554
lambda: 1 | epoch: 100 | error: 0.4554
lambda: 10 | epoch: 100 | error: 0.4554
lambda: 100 | epoch: 100 | error: 0.4554
lambda: 1000 | epoch: 100 | error: 0.4554
lambda: 0.001 | epoch: 500 | error: 0.4996
lambda: 0.01 | epoch: 500 | error: 0.45213333333333333
lambda: 0.1 | epoch: 500 | error: 0.4554
lambda: 1 | epoch: 500 | error: 0.4554
lambda: 10 | epoch: 500 | error: 0.46026666666666666
lambda: 100 | epoch: 500 | error: 0.4554
lambda: 1000 | epoch: 500 | error: 0.4554

```

```


[79]: plt.subplot(121)
      plt.plot(epochs * len(lambda_s), error, 'go')
      plt.xlabel("epoch")
      plt.subplot(122)
      plt.plot(lambda_s * len(epochs), error, 'ro')
      plt.xlabel("lambda")

```

```

[79]: Text(0.5, 0, 'lambda')


```



```
[0]: # Computing the prediction error on the testing set
```

```
[80]: coef2 = perform_sgd(train_scores[:5000], train_label[:5000], 0.01, 0.01, 500,
    ↪ decrease = True, plot = True)
    error2 = compute_mis(coef2, test_scores, test_label)
    print("test error: ", error2)
```

```
test error: 0.4534
```



[0]: *# In general, the previous model to which we fed overall reviews is more
effective because the score models have generated higher error rates. The
results are acceptable considering the fact that texts provide more_
→information
about ratings than mere numbers.*

[0]:

[0]:

[0]:

[0]: